
Stock Price Forecasting using LSTM

May 31, 2021

Leonardo Emili Alessio Luciani

1. Introduction

Stock price prediction is an active field of study where the objective goal is to predict the trend of the stock market, typically based on its historical evidence. Over the years, many methods have been proposed to tackle the task: ranging from the Naive Forecast approach, which trivially forecasts the stock value to be the last observed one, up to the most recent ones that heavily rely on machine learning techniques. In this context, we will delve into the topic of stock price prediction using fundamental data to assess whether a stock is attractive to investors. This technique, as opposed to technical analysis, considers corporate statics, as well as financial reports, to model the problem. Here, the underlying assumption is that one cannot tell whether it's worth buying a stock only by looking at its current price and volumes. Indeed, quarterly released financial reports about companies and external information (e.g. the common sentiment on a stock ticker) may be crucial to assess the quality of a given stock. In this project, we experimented with the effectiveness of deep learning techniques applied to the stock price prediction task. The key idea is the use of recurrent neural networks that are able to exploit temporal dependencies of events, hence conditioning the presence of an event at timestep t on the previous events at $t-1, t-2, \dots$.

2. Related work

Some approaches have been explored in this direction, such as in (Zou & Qu, 2020) where the authors apply LSTM, Stacked-LSTM, and Attention-Based LSTM into the prediction of stock prices. The authors also propose an evaluating framework to assess the quality of their models based on the return of the trading strategy. In another work (Mehtab et al., 2020), the authors predict the open value of NIFTY 50 using different machine learning and deep learning models. They also demonstrate that using one-

week prior data as input leads to good results.

3. Dataset

Data wrangling operations have been crucial for this task. In fact, we start considering the S&P 500 stock data, which is a collection of daily stock prices for all companies from the S&P 500 index. This dataset provides us with the following features: open price, highest price, lowest price, close price, volume, stock ticker, and date. We also consider an auxiliary dataset that provides us with fundamental data from Yahoo Finance. The dataset contains the following columns: Forward P/E, DE Ratio, Earnings Growth, Enterprise Value/EBITDA, EBITDA, Current Ratio, Cash Flow, Trailing P/E, Beta, PEG Ratio, Gross Profit, Total Debt, Price, Return on Equity, Return on Assets, Price/Book, Revenue Growth, Operating Margin, Enterprise Value/Revenue, Revenue, Total Cash, Enterprise Value, Total Cash Per Share, Profit Margin, Price/Sales, Book Value Per Share, Diluted EPS, Market Cap, Revenue Per Share, Net Income Avl to Common, Ticker, Date. Based on the date column, we can align the two datasets such that for each event in the S&P 500 dataset, we have the latest available financial report. From now on, we will refer to the dataset obtained from the alignment process simply as the dataset.

3.1. Feature engineering

We fast-forward values whenever possible, otherwise replacing missing values with constant values (i.e. zero paddings). The intuition here is that if there are missing values at the beginning of the history of a stock, it means that the feature is not available, otherwise referring to the latest value as the most updated one. As an example, consider holidays when markets are closed, and the price does not change since no orders are placed. We also fill the dataset with missing working days using the same fast-forward strategy. Furthermore, we perform massive feature engineering steps adding technical indicators such as SMA and RSI. (**TODO:** Explain how these indicators work and why we think they are important for our task, explain feature scaling, explain 2012-2013 train/dev/test splits). In order to predict the target adjusted close price, we consider

Email: Leonardo Emili
<emili.1802989@studenti.uniroma1.it>, Alessio Luciani
<luciani.1797637@studenti.uniroma1.it>.

the dataset according to the sliding window approach, also known as the lag method. In this way, we decompose the original dataset with overlapping windows and condition the target value only on its lag. During the training phase, we treat both the step size (i.e. the number of days before a new window) and the window size (i.e. the size of the lag) as hyperparameters and tune them accordingly.

4. Models

In this project, we apply the sliding window approach with deep recurrent neural networks. As our first models, we start with naive recurrent networks, respectively LSTM and GRU models. As opposed to vanilla RNN, LSTM and GRU networks partially solve the problem of vanishing gradient by employing a gating mechanism to regulate the amount of information to carry from previous time steps. Our input has the following shape: (*batch_size*, *window_size*, *feature_size*). Moreover, we assume that there exists a local pattern that we can leverage to predict future prices. In fact, in a third model, we make use of CNN layers to extract such information across time steps. It consists of a convolutional layer followed by an LSTM layer, then connected to two dense layers. As a fourth model, we employ the attention mechanism over our input data and then feed it to an LSTM layer. In particular, we employ a multi-head self-attention mechanism *to jointly attend to information from different representation subspaces at different positions* (Vaswani et al., 2017).

5. Hyperparameters

In this section, we present a subset of the hyperparameters used. However, it is not intended to be an exhaustive list of all the hyperparameter. The list of runs is logged using wandb (Biewald, 2020) and can be reached at [link](#).

Table 1. List of hyperparameters.

Hyperparameter	Value
Window size	20
Step size	1
Optimizer	SGD
Batch size	1024

6. Experimental results

As the evaluation framework, we consider multiple metrics in order to assess the quality of a particular model. Since we are dealing with a regression task it comes naturally to adopt the Mean Squared Error (MSE) measure as a proxy to indicate how well a model performs. We also consider the R-squared statistical measure to check how similar the predicted and the observed data are similar. However, since

we are aware that R-squared increases when adding more independent variables, we use the adjusted r-squared. Its formulation is similar to the one of r-square and only differs from the degrees of freedom. Finally, we measure the real performances of our system according to the return of a trading strategy. The system is designed to buy whenever the predicted price is higher than the current one, and when the day after we close the position, we register the gain or the loss. **TODO:** explain what are our best performing models and the reasons why they perform so well.

Table 2. Performance comparison.

	MSE	R2	Adjusted R2	Operation accuracy (%)	Return (%)
Naive LSTM	0.045	0.985	0.955	85.79	419.12
Naive GRU	0.062	0.988	0.970	85.11	422.87
CNN-LSTM	0.033	0.988	0.968	83.52	390.87
Attention-based LSTM	0.022	0.991	0.977	0.8102	420.26

7. Conclusions

TODO: Add conclusions here ..

References

- Biewald, L. Experiment tracking with weights and biases, 2020. URL <https://www.wandb.com/>. Software available from wandb.com.
- Mehtab, S., Sen, J., and Dutta, A. Stock price prediction using machine learning and lstm-based deep learning models, 2020.
- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, L., and Polosukhin, I. Attention is all you need. 2017. URL <https://arxiv.org/pdf/1706.03762.pdf>.
- Zou, Z. and Qu, Z. Using lstm in stock prediction and quantitative trading, 2020.